

```
In [1]: import numpy as np
import torch
import cv2
```

```
In [2]: from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
from torch.utils.data import DataLoader
from torch import nn
import matplotlib.pyplot as plt
```

```
In [3]: # Get cpu, gpu or mps device for training.
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")
```

Using cpu device

```
In [4]: #You may need this to access the datasets.
import ssl
ssl._create_default_https_context = ssl._create_unverified_context
```

```
In [5]: #This is the training data.
from torchvision.datasets import CIFAR10
train_data = datasets.CIFAR10(root="train/",
train=True,
download=True,
transform = ToTensor())
```

Files already downloaded and verified

```
In [6]: #Test data.
test_data = datasets.CIFAR10(root="test/",
train=False,
download=True,
transform = ToTensor())
```

Files already downloaded and verified

```
In [7]: print(train_data)
```

```
Dataset CIFAR10
  Number of datapoints: 50000
  Root location: train/
  Split: Train
  StandardTransform
  Transform: ToTensor()
```

```
In [8]: print(test_data)
```

```
Dataset CIFAR10
  Number of datapoints: 10000
  Root location: test/
  Split: Test
  StandardTransform
  Transform: ToTensor()
```

```
In [9]: dir(train_data)
```

```
Out[9]: ['__add__',
          '__annotations__',
          '__class__',
          '__class_getitem__',
          '__delattr__',
          '__dict__',
          '__dir__',
          '__doc__',
          '__eq__',
          '__format__',
          '__ge__',
          '__getattr__',
          '__getitem__',
          '__getstate__',
          '__gt__',
          '__hash__',
          '__init__',
          '__init_subclass__',
          '__le__',
          '__len__',
          '__lt__',
          '__module__',
          '__ne__',
          '__new__',
          '__orig_bases__',
          '__parameters__',
          '__reduce__',
          '__reduce_ex__',
          '__repr__',
          '__setattr__',
          '__sizeof__',
          '__slots__',
          '__str__',
          '__subclasshook__',
          '__weakref__',
          '_check_integrity',
          '_format_transform_repr',
          '_is_protocol',
          '_load_meta',
          '_repr_indent',
          'base_folder',
          'class_to_idx',
          'classes',
          'data',
          'download',
          'extra_repr',
          'filename',
          'meta',
          'root',
          'target_transform',
          'targets',
          'test_list',
          'tgz_md5',
          'train',
          'train_list',
          'transform',
          'transforms',
          'url']
```

```
In [10]: train_data.class_to_idx
```

```
Out[10]: {'airplane': 0,
          'automobile': 1,
          'bird': 2,
          'cat': 3,
          'deer': 4,
          'dog': 5,
          'frog': 6,
          'horse': 7,
          'ship': 8,
          'truck': 9}
```

```
In [11]: test_data.class_to_idx
```

```
Out[11]: {'airplane': 0,
          'automobile': 1,
          'bird': 2,
          'cat': 3,
          'deer': 4,
          'dog': 5,
          'frog': 6,
          'horse': 7,
          'ship': 8,
          'truck': 9}
```

```
In [12]: batch_size = 64
```

```
# Create data loaders.
train_dataloader = DataLoader(train_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)

for X, y in train_dataloader:
    print(f"Shape of X [N, C, H, W]: {X.shape}")
    print(f"Shape of y: {y.shape} {y.dtype}")
    break
```

```
Shape of X [N, C, H, W]: torch.Size([64, 3, 32, 32])
Shape of y: torch.Size([64]) torch.int64
```

```
In [13]: def label_array(dataloader):
          # initialize the empty list
          cifar10_test_labels = []

          # read data from dataloader
          for inputs, labels in dataloader:
              # add labels to list cifar10_test_labels
              cifar10_test_labels.append(labels)

          # convert cifar10_test_labels to tensor
          cifar10_test_labels = torch.cat(cifar10_test_labels, dim=0)

          return cifar10_test_labels
```

```
In [14]: # check function label_array
          label_array(test_dataloader)
```

```
Out[14]: tensor([3, 8, 8, ..., 5, 1, 7])
```

```
In [15]: label_array(test_dataloader).shape
```

```
Out[15]: torch.Size([10000])
```

```
In [16]: class CIFARConvNet(nn.Module):
    def __init__(self):
        super(CIFARConvNet, self).__init__()
        #the convolution layers
        self.conv1 = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding='sa
            nn.ReLU()
        )
        self.conv2 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=5, padding='s
            nn.ReLU()
        )
        #the MaxPool layers
        self.pool = nn.MaxPool2d(kernel_size=2)
        nn.Dropout(0.25)

        #the standard layers
        self.fc1 = nn.Sequential(
            nn.Flatten(),
            nn.Linear(4096, 256),
            nn.ReLU(),
            nn.Dropout(0.4)
        )
        self.fc2 = nn.Sequential(
            nn.Linear(256,10)
        )

    def forward(self, x):
        x = self.conv1(x)
        x = self.pool(x)
        x = self.conv2(x)
        x = self.pool(x)
        x = self.fc1(x)
        x = self.fc2(x)

    return x
```

```
In [148... # declare variable to store training history
train_his = {
    "train_loss": [],
    "train_acc": [],
    "val_loss": [],
    "val_acc": []
}
# split train_data into train_subset and validation_subset
train_subset, val_subset = torch.utils.data.random_split(
    train_data, [40000, 10000], generator=torch.Generator().manual_seed(1))

# create dataloader for train_subset and validation_subset
train_dataloader = DataLoader(train_subset, batch_size=batch_size)
val_dataloader = DataLoader(val_subset, batch_size=batch_size)

# calculate steps per epoch for training set
```

```
train_steps = len(train_dataloader.dataset) // batch_size
val_steps = len(val_dataloader.dataset) // batch_size
```

In [159...

```
import time
lr = 1e-3
num_epochs = 60
# initialize model
model = CIFARConvNet().to(device)
# select loss function
loss_fn = nn.CrossEntropyLoss()
# select optimizer
optimizer = torch.optim.Adam(model.parameters(), lr=lr)

# measure how long training is going to take
print("[INFO] training the network...")
startTime = time.time()

for epochs in range(num_epochs):
    # set the model in training mode
    model.train()
    # initialize the total training and validation loss
    total_train_loss = 0
    total_val_loss = 0
    # initialize the number of correct predictions in the training step
    train_accuracy = 0
    val_accuracy = 0
    # loop over the training set
    for (x, y) in train_dataloader:
        # send the input to the device
        (x, y) = (x.to(device), y.to(device))
        # perform a forward pass and calculate the training loss
        pred = model(x)
        loss = loss_fn(pred, y)
        # zero out the gradients, perform the backpropagation step, and update t
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        # add the loss to the total training loss so far and calculate the numbe
        total_train_loss += loss
        train_accuracy += (pred.argmax(1) == y).type(torch.float).sum().item()

    # switch off autograd for evaluation
    with torch.no_grad():
        # set the model in evaluation mode
        model.eval()
        # loop over the validation set
        for (x, y) in val_dataloader:
            # send the input to the device
            (x, y) = (x.to(device), y.to(device))
            # make the predictions and calculate the validation loss
            pred = model(x)
            total_val_loss += loss_fn(pred, y)
            # calculate the number of correct predictions
            val_accuracy += (pred.argmax(1) == y).type(torch.float).sum().item()

    # calculate the average training loss and validation loss
    avg_train_loss = total_train_loss / train_steps
    avg_val_loss = total_val_loss / val_steps

    # calculate the training accuracy
```

```
train_accuracy = train_accuracy / len(train_dataloader.dataset)
val_accuracy = val_accuracy / len(val_dataloader.dataset)

# update our training history
train_his["train_loss"].append(avg_train_loss.cpu().detach().numpy())
train_his["train_acc"].append(train_accuracy)
train_his["val_loss"].append(avg_val_loss.cpu().detach().numpy())
train_his["val_acc"].append(val_accuracy)

# print the model training information
print("    EPOCH: {}/{}".format(epochs + 1, num_epochs))
print("    Train loss: {:.6f}, Train accuracy: {:.4f}".format(avg_train_l
print("    Validation loss: {:.6f}, Validation accuracy: {:.4f}".format(a

# finish measuring how long training took
endTime = time.time()
print("[INFO] total time taken to train the model: {:.2f}s".format(endTime - sta
```

[INFO] training the network...

EPOCH: 1/60
Train loss: 1.585631, Train accuracy: 0.4235
Validation loss: 1.295445, Validation accuracy: 0.5347

EPOCH: 2/60
Train loss: 1.240099, Train accuracy: 0.5583
Validation loss: 1.102593, Validation accuracy: 0.6119

EPOCH: 3/60
Train loss: 1.081914, Train accuracy: 0.6166
Validation loss: 1.014575, Validation accuracy: 0.6363

EPOCH: 4/60
Train loss: 0.984198, Train accuracy: 0.6538
Validation loss: 1.001759, Validation accuracy: 0.6458

EPOCH: 5/60
Train loss: 0.903931, Train accuracy: 0.6805
Validation loss: 0.959133, Validation accuracy: 0.6648

EPOCH: 6/60
Train loss: 0.841330, Train accuracy: 0.7029
Validation loss: 0.952182, Validation accuracy: 0.6652

EPOCH: 7/60
Train loss: 0.786465, Train accuracy: 0.7238
Validation loss: 0.898338, Validation accuracy: 0.6879

EPOCH: 8/60
Train loss: 0.724433, Train accuracy: 0.7420
Validation loss: 0.921266, Validation accuracy: 0.6811

EPOCH: 9/60
Train loss: 0.678609, Train accuracy: 0.7597
Validation loss: 0.889362, Validation accuracy: 0.6898

EPOCH: 10/60
Train loss: 0.635303, Train accuracy: 0.7738
Validation loss: 0.907758, Validation accuracy: 0.6935

EPOCH: 11/60
Train loss: 0.595771, Train accuracy: 0.7856
Validation loss: 0.913831, Validation accuracy: 0.6873

EPOCH: 12/60
Train loss: 0.564334, Train accuracy: 0.7983
Validation loss: 0.935990, Validation accuracy: 0.6939

EPOCH: 13/60
Train loss: 0.529436, Train accuracy: 0.8104
Validation loss: 0.981988, Validation accuracy: 0.6916

EPOCH: 14/60
Train loss: 0.498014, Train accuracy: 0.8206
Validation loss: 0.982573, Validation accuracy: 0.6915

EPOCH: 15/60
Train loss: 0.465997, Train accuracy: 0.8311
Validation loss: 1.044847, Validation accuracy: 0.6920

EPOCH: 16/60
Train loss: 0.429090, Train accuracy: 0.8447
Validation loss: 1.082340, Validation accuracy: 0.6952

EPOCH: 17/60
Train loss: 0.412586, Train accuracy: 0.8483
Validation loss: 1.088463, Validation accuracy: 0.6861

EPOCH: 18/60
Train loss: 0.381381, Train accuracy: 0.8599
Validation loss: 1.142104, Validation accuracy: 0.6900

EPOCH: 19/60
Train loss: 0.374458, Train accuracy: 0.8629
Validation loss: 1.114530, Validation accuracy: 0.6952

EPOCH: 20/60
Train loss: 0.350023, Train accuracy: 0.8735

Validation loss: 1.198752, Validation accuracy: 0.6846
EPOCH: 21/60
Train loss: 0.340961, Train accuracy: 0.8742
Validation loss: 1.254526, Validation accuracy: 0.6850
EPOCH: 22/60
Train loss: 0.317285, Train accuracy: 0.8834
Validation loss: 1.321228, Validation accuracy: 0.6758
EPOCH: 23/60
Train loss: 0.305987, Train accuracy: 0.8886
Validation loss: 1.305097, Validation accuracy: 0.6871
EPOCH: 24/60
Train loss: 0.282666, Train accuracy: 0.8949
Validation loss: 1.324823, Validation accuracy: 0.6830
EPOCH: 25/60
Train loss: 0.285432, Train accuracy: 0.8963
Validation loss: 1.413260, Validation accuracy: 0.6815
EPOCH: 26/60
Train loss: 0.274008, Train accuracy: 0.8981
Validation loss: 1.441490, Validation accuracy: 0.6810
EPOCH: 27/60
Train loss: 0.262669, Train accuracy: 0.9039
Validation loss: 1.424560, Validation accuracy: 0.6858
EPOCH: 28/60
Train loss: 0.251488, Train accuracy: 0.9083
Validation loss: 1.477663, Validation accuracy: 0.6867
EPOCH: 29/60
Train loss: 0.243733, Train accuracy: 0.9092
Validation loss: 1.501709, Validation accuracy: 0.6883
EPOCH: 30/60
Train loss: 0.239148, Train accuracy: 0.9133
Validation loss: 1.486093, Validation accuracy: 0.6926
EPOCH: 31/60
Train loss: 0.219644, Train accuracy: 0.9192
Validation loss: 1.579766, Validation accuracy: 0.6856
EPOCH: 32/60
Train loss: 0.217019, Train accuracy: 0.9211
Validation loss: 1.509071, Validation accuracy: 0.6959
EPOCH: 33/60
Train loss: 0.216126, Train accuracy: 0.9196
Validation loss: 1.567184, Validation accuracy: 0.6887
EPOCH: 34/60
Train loss: 0.208557, Train accuracy: 0.9246
Validation loss: 1.655564, Validation accuracy: 0.6891
EPOCH: 35/60
Train loss: 0.201002, Train accuracy: 0.9269
Validation loss: 1.628972, Validation accuracy: 0.6922
EPOCH: 36/60
Train loss: 0.201495, Train accuracy: 0.9270
Validation loss: 1.661271, Validation accuracy: 0.6884
EPOCH: 37/60
Train loss: 0.200366, Train accuracy: 0.9270
Validation loss: 1.659268, Validation accuracy: 0.6832
EPOCH: 38/60
Train loss: 0.191297, Train accuracy: 0.9311
Validation loss: 1.627513, Validation accuracy: 0.6899
EPOCH: 39/60
Train loss: 0.179723, Train accuracy: 0.9335
Validation loss: 1.706143, Validation accuracy: 0.6863
EPOCH: 40/60
Train loss: 0.177106, Train accuracy: 0.9358

Validation loss: 1.770249, Validation accuracy: 0.6869
EPOCH: 41/60
Train loss: 0.181447, Train accuracy: 0.9335
Validation loss: 1.719392, Validation accuracy: 0.6875
EPOCH: 42/60
Train loss: 0.177426, Train accuracy: 0.9358
Validation loss: 1.767121, Validation accuracy: 0.6899
EPOCH: 43/60
Train loss: 0.169567, Train accuracy: 0.9395
Validation loss: 1.877388, Validation accuracy: 0.6889
EPOCH: 44/60
Train loss: 0.163309, Train accuracy: 0.9399
Validation loss: 1.823595, Validation accuracy: 0.6879
EPOCH: 45/60
Train loss: 0.168265, Train accuracy: 0.9389
Validation loss: 1.905448, Validation accuracy: 0.6838
EPOCH: 46/60
Train loss: 0.166925, Train accuracy: 0.9393
Validation loss: 1.871604, Validation accuracy: 0.6821
EPOCH: 47/60
Train loss: 0.171206, Train accuracy: 0.9378
Validation loss: 1.880969, Validation accuracy: 0.6878
EPOCH: 48/60
Train loss: 0.158327, Train accuracy: 0.9422
Validation loss: 1.915547, Validation accuracy: 0.6870
EPOCH: 49/60
Train loss: 0.157527, Train accuracy: 0.9427
Validation loss: 1.867690, Validation accuracy: 0.6864
EPOCH: 50/60
Train loss: 0.157031, Train accuracy: 0.9434
Validation loss: 1.969535, Validation accuracy: 0.6922
EPOCH: 51/60
Train loss: 0.148118, Train accuracy: 0.9478
Validation loss: 1.985334, Validation accuracy: 0.6918
EPOCH: 52/60
Train loss: 0.149415, Train accuracy: 0.9441
Validation loss: 1.946991, Validation accuracy: 0.6933
EPOCH: 53/60
Train loss: 0.142246, Train accuracy: 0.9475
Validation loss: 1.949271, Validation accuracy: 0.6886
EPOCH: 54/60
Train loss: 0.144950, Train accuracy: 0.9482
Validation loss: 1.992840, Validation accuracy: 0.6889
EPOCH: 55/60
Train loss: 0.147576, Train accuracy: 0.9472
Validation loss: 2.061358, Validation accuracy: 0.6958
EPOCH: 56/60
Train loss: 0.145506, Train accuracy: 0.9487
Validation loss: 1.994285, Validation accuracy: 0.6895
EPOCH: 57/60
Train loss: 0.141268, Train accuracy: 0.9499
Validation loss: 1.984730, Validation accuracy: 0.6892
EPOCH: 58/60
Train loss: 0.139257, Train accuracy: 0.9507
Validation loss: 2.011512, Validation accuracy: 0.6935
EPOCH: 59/60
Train loss: 0.140522, Train accuracy: 0.9498
Validation loss: 1.989849, Validation accuracy: 0.6929
EPOCH: 60/60
Train loss: 0.134466, Train accuracy: 0.9516

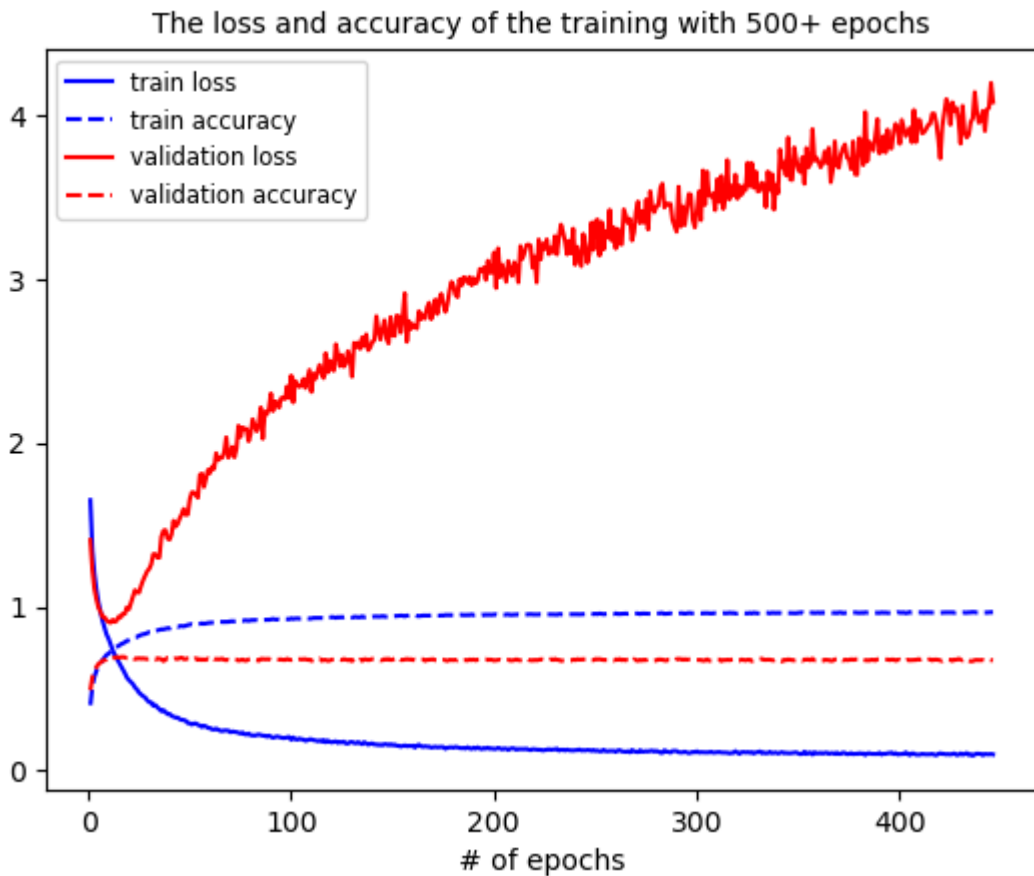
Validation loss: 2.142626, Validation accuracy: 0.6915
 [INFO] total time taken to train the model: 2701.16s

```
In [158... # this variable store the result of training with 400+ epochs
len(train_his_400plus["val_acc"])
```

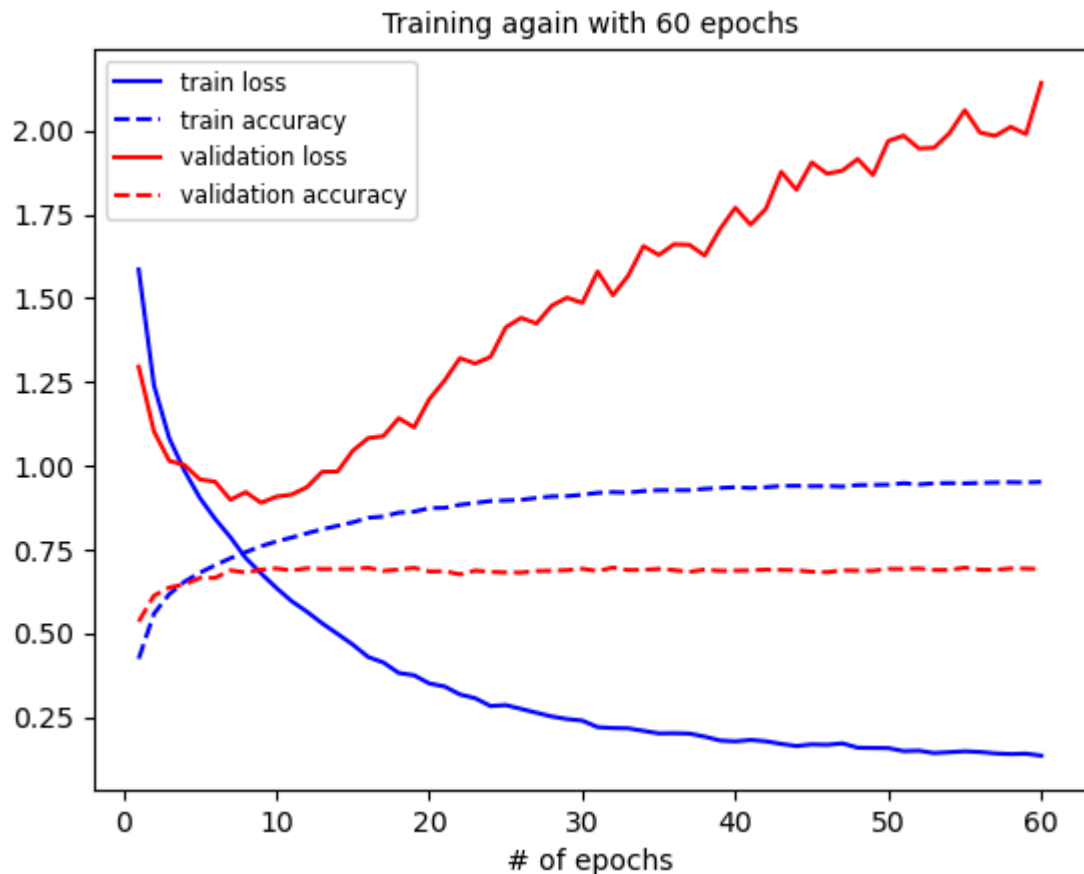
Out[158... 446

```
In [142... # function for plotting the loss and accuracy
def plot_train_loss(loss_dict, txt):
    # create data for y-axis
    y1=np.array(loss_dict["train_loss"])
    y2=np.array(loss_dict["train_acc"])
    y3=np.array(loss_dict["val_loss"])
    y4=np.array(loss_dict["val_acc"])
    # create data for x-axis
    X = np.linspace(1, len(y1), num=len(y1)).reshape(-1,1)
    # plotting
    plt.xlabel('# of epochs')
    plt.plot(X, y1, 'b-')
    plt.plot(X, y2, 'b--')
    plt.plot(X, y3, 'r-')
    plt.plot(X, y4, 'r--')
    plt.legend(['train loss', 'train accuracy', 'validation loss', 'validation accu
    plt.title(txt, fontsize=10)
```

```
In [157... plot_train_loss(train_his_400plus, "The loss and accuracy of the training with 5
```



```
In [160... plot_train_loss(train_his, "Training again with 60 epochs")
```



```
In [161... # save trained model to local directory
torch.save(model.state_dict(), "cifar_convnet1")
```

```
In [162... # Load the model
model1 = CIFARConvNet().to(device)
model1.load_state_dict(torch.load("cifar_convnet1", weights_only=True))
```

```
Out[162... <All keys matched successfully>
```

```
In [163... def get_prediction(dataloader, model):
    pred = []
    # set model to prediction state
    model.eval()
    # Loop through the dataloader
    for inputs, labels in dataloader:
        with torch.no_grad():
            # send the input to the device
            inputs = inputs.to(device)
            # make the predictions
            outputs = model(inputs)
            # identify the predicted class
            predict = outputs.max(dim=1)
            # add them to the list
            pred.append(predict.indices)

    # convert list to tensor
    pred = torch.cat(pred, dim=0)
    return pred
```

```
In [164... def calculate_accuracy(test_label, predicted):
    # calculate the accuracy
```

```
acc = sum(test_label==predicted).item()/len(test_label)
return acc
```

```
In [165... def confusion_dict(test_label, predicted):
    conf_dict = {}
    # convert tensor to numpy
    test_np = test_label.numpy()
    predicted_np = predicted.numpy()

    for key in range(10):
        # filter the incorrect prediction in predicted list
        error_label = predicted_np[(test_np==key) & (predicted_np!=key)]
        # calculate the frequency of each incorrect label
        unique, counts = np.unique(error_label, return_counts=True)
        # add the top 2 incorrect predictions into conf_dict
        if len(counts) < 2:
            conf_dict[key] = tuple(unique)
        else:
            conf_dict[key] = tuple(unique[np.argsort(counts)[len(counts)-2:]]

    return conf_dict
```

```
In [166... # predict
predicted = get_prediction(test_dataloader, model1)
```

```
In [167... # get test label
test_label = label_array(test_dataloader)
```

```
In [168... # accuracy
calculate_accuracy(test_label, predicted)
```

```
Out[168... 0.6897
```

```
In [169... # confusion dictionary
conf_dict = confusion_dict(test_label, predicted)
conf_dict
```

```
Out[169... {0: (2, 8),
1: (8, 9),
2: (4, 5),
3: (6, 5),
4: (2, 7),
5: (2, 3),
6: (2, 3),
7: (4, 5),
8: (1, 0),
9: (0, 1)}
```

```
In [170... # function for giving comments based on the confusion dictionary
def comment(conf_dict):
    # get class based on index
    def get_class(index):
        return list(test_data.class_to_idx.keys())[list(test_data.class_to_idx.v

    # plot first incorrect prediction
    def plot_top_diff(test_label, predicted, txt, key1, key2, key3=-1):
        test_np = test_label.numpy()
        predicted_np = predicted.numpy()
        # get index of images with label key1
```

```

key1_idx = np.argwhere(test_np==key1)
# filter prediction of image with label key1
key1_predict = predicted_np[test_np==key1]
# get index of prediction: key1 to key2
key2_idx = np.argwhere(key1_predict==key2)
# plot first incorrect predicted image: key1 to key2
plt.figure(figsize=(5, 2))
plt.figtext(0.5, -0.1, txt, wrap=True, horizontalalignment='center', font
plt.subplot(121); plt.imshow(test_data.data[key1_idx[key2_idx[0][0]][0])
plt.title("{} >> {}".format(get_class(key1),get_class(key2)), fontsize=1
if key3 > -1:
    # get index of prediction: key1 to key3
    key3_idx = np.argwhere(key1_predict==key3)
    # plot first incorrect predicted image: key1 to key3
    plt.subplot(122); plt.imshow(test_data.data[key1_idx[key3_idx[0][0]][0])
    plt.title("{} >> {}".format(get_class(key1),get_class(key3)), fontsi

# print comment
for key in conf_dict.keys():
    no_of_err = 1 if type(conf_dict[key]) is int else len(conf_dict[key])
    if no_of_err==0:
        print("{} are classified properly.\n".format(get_class(key)))
    elif no_of_err==1:
        err = 0 if type(conf_dict[key]) is int else conf_dict[key][0]
        txt = "{} are confused with {}.\n".format(get_class(key),
                                                    get_class(err))
        plot_top_diff(test_label, predicted, txt, key, err)
    else:
        txt = "{} are confused with {} and {}.\n".format(get_class(key),
                                                            get_class(conf_dict
                                                            get_class(conf_dict
        plot_top_diff(test_label, predicted, txt, key, conf_dict[key][0], co

```

In [171... *# comment on the errors*
comment(conf_dict)

airplane >> bird



airplane >> ship



airplane are confused with bird and ship.

automobile >> ship



automobile >> truck



automobile are confused with ship and truck.

bird >> deer



bird >> dog




bird are confused with deer and dog.

cat >> frog



cat >> dog



cat are confused with frog and dog.

deer >> bird




deer >> horse



deer are confused with bird and horse.

dog >> bird



dog >> cat



dog are confused with bird and cat.

frog >> bird




frog >> cat



frog are confused with bird and cat.

horse >> deer



horse >> dog




horse are confused with deer and dog.

ship >> automobile



ship >> airplane



ship are confused with automobile and airplane.

truck >> airplane



truck >> automobile



truck are confused with airplane and automobile.